

GPGPU programming

General-purpose Processing on Graphics Processing Units

Robin Faury
robinfaurypro@gmail.com

12-14-2021

Introduction

- ▶ The purpose of parallel processing
- ▶ What is a graphic card?
- ▶ Languages
- ▶ GPGPU usage in the industry
- ▶ Q&A

The purpose of parallel processing

Moore's Law

Every two years, the density of transistors in an integrated circuit doubles. That means we can compute the critical path of an algorithm faster.



To infinity and beyond!

Critical path

Sometimes, algorithms process data one by one. When applicable, it is necessary to find the critical path and execute it in parallel. Modern CPUs offer the ability to run some threads at the same time. However, CPUs don't have a lot of cores available. For massive parallel computation we will use GPUs.



Latency and Throughput

- ▶ **Latency:** This is the time between an action and the response to this action. For example, a key press event and its process.
- ▶ **Throughput:** This is the rate of production. For example, the number of pixels processed in one second.

CPU versus GPU

- ▶ The aim of **CPU** is to be very responsive. For that they adopt strategies to hide the latency (pre-fetch, branch prediction...). However, these algorithms need a lot of memory cache.
- ▶ The aim of **GPU** is to process a lot of data. That why there is way more cores into GPU than CPU. The memory cache need to be divided by the number of cores.

A CPU has a small amount of thread, so it can dedicate a large cache to store intermediate result to be very responsive. A CPU has a low latency but a low throughput too.

A GPU need a setup before running a process and he cannot store a lot of intermediate results but it can run a very large amount of thread. A GPU has a high latency but a high throughput too.

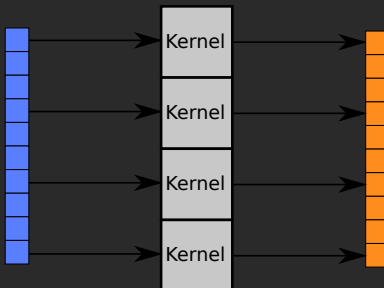
Flynn's taxonomy

We just see two architectures to process data (the CPU and the GPU) but we can imagine other ones. Flynn made a classification of computer architectures.

- ▶ **SISD**: Single Instruction Single Data. One core of a CPU.
- ▶ **SIMD**: Single Instruction Multiple Data. A GPU.
- ▶ **MISD**: Multiple Instruction Single Data. An architecture that need to compute redundancy on a system.
- ▶ **MIMD**: Multiple Instruction Multiple Data. A distributed system.

SIMD: a world of buffers

General Purpose computing on GPU leverage the power of GPU to process a stream of data. All data will be store into a buffer and the solving of the arithmetic computation will be done from this buffer to another one (we can read/write a buffer but only in some specific cases). One process is called a kernel for the GPGPU or a shader for the graphics pipeline.

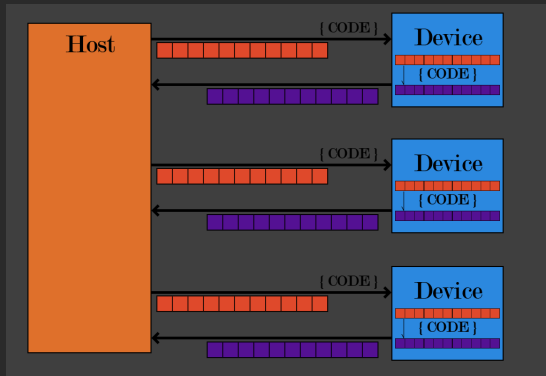


Host and Devices

- ▶ Host: The CPU and its memory. The host can manage the memory on both the host and the device. The executed code can launch kernels.
- ▶ Devices: The GPU and its memory. Kernels are executed on many GPU threads in parallel.

Host and Device

A GPGPU need to be seen as an asynchronous service. The host (CPU) will send to a device (GPU) the assembly code of the process and then the data to process. The device don't need anymore information to work and the host is free to run other task. You have one host and one or more devices.



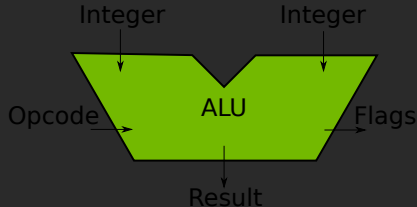
What is a graphic card?

History

The first Graphics Processing Unit (GPU) was used for drawing game sprites. It was a dedicated device for formatted data. Ten years after we had the ability to draw lines, fill areas and control the blitter. In 1990, the graphical API appears and allows us to send assembly code to the device.

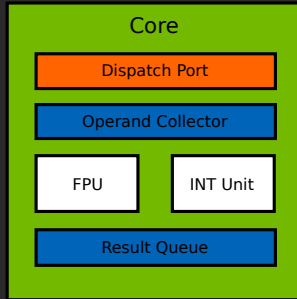
Arithmetic Logic Unit

The Arithmetic Logic Unit (ALU) is the component that performs arithmetic operations. The GPU is more focused on floating point operations, multiple ALUs are combined to create a Floating Point Unit (FPU). Latest GPU include also dedicated chip for double operations.



Core

Cores are used to execute opcodes from compiled kernels. They are composed of an FPU, logic unit, branch unit and compare unit.



Streaming Multiprocessor

The Streaming Multiprocessor (SM) organizes threads in groups of 32 called warp.



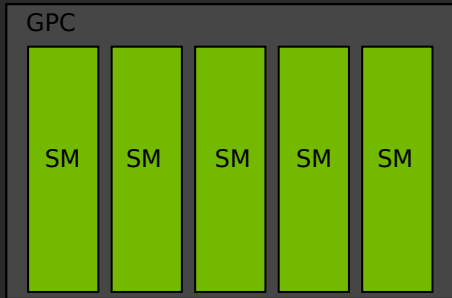
Streaming Multiprocessor

On the GP104 (The GPU of GTX 1080) each SM has four warps.



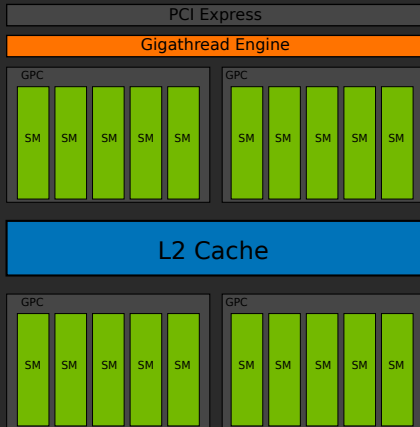
Graphics Processing Clusters

A Graphics Processing Clusters (GPC) is a collection of streaming multiprocessors. In the case of the GP104, there are four clusters.



GP104

All the GPC are connected to the L2 cache memory. The Gigathread engine distributes block threads to streaming multiprocessor. This device has $32 \text{ cores} * 4 \text{ warps} * 5 \text{ SMs} * 4 \text{ GPCs} = 2560 \text{ CUDA cores}$.



GA102



GPGPU languages

APIs

To communicate with the graphic card we need API. Each one have its pro and con.

- ▶ **CUDA:** The Nvidia API. Very close to the C++, easy to use. Works only on Nvidia card.
- ▶ **DirectX:** The Windows API. Design for game development on Windows.
- ▶ **Metal:** The Apple API. Design for graphic development on MacOS.
- ▶ **OpenGL:** Developed by Khronos, design for all OS. Deprecated from 2019.
- ▶ **Vulkan:** Developed by Khronos, design for all OS but MacOS. Launched on 2018. Difficult to use.

Dispatch threads

From the host, you can dispatch a large amount of threads. But those threads need to be organized. On a kernel we specify the number of invocations and when we use the dispatch function we specify the number of working groups. Each working group will run x instances of the kernel.

If we specify 42 invocations in the kernel code and we dispatch 100 working groups, the GPU will run the kernel 42 000 times.

Kernel - GLSL

The kernel is the code that GPU will execute on the data. First it need the version of the GLSL language. It also need a main function.

```
#version 430 core
```

```
void main() {
```

```
}
```


Kernel - GLSL

We can specify the number of invocation of the kernel. As you can see, the invocation can be set in 1D, 2D or 3D. It's useful if you work on buffer, texture or volume.

```
#version 430 core

layout(
    local_size_x = 42,
    local_size_y = 1,
    local_size_z = 1) in;

void main() {

}
```

Kernel - GLSL

If you need a input or an output buffer, you need to specify the following layout:

```
#version 430 core
```

```
layout(  
    local_size_x = 42,  
    local_size_y = 1,  
    local_size_z = 1) in;
```

```
layout(std140, binding = 0) buffer MyBuffer {  
    float buffer_[];  
};
```

```
void main() {  
  
}
```

Kernel - GLSL

And finally, we can read/write this buffer.

```
#version 430 core
```

```
layout(  
    local_size_x = 42,  
    local_size_y = 1,  
    local_size_z = 1) in;
```

```
layout(std140, binding = 0) buffer MyBuffer {  
    float buffer_[];  
};
```

```
void main() {  
    float data=buffer[gl_GlobalInvocationID.x].x;  
    ++data;  
    buffer[gl_GlobalInvocationID.x]=data;  
}
```

Indexing

In the last example, we use the `GlobalInvocationID` variable as the index of our buffer. This value can be computed using the $\text{WorkGroupID} * \text{WorkGroupSize} + \text{LocalInvocationID}$;

Indexing with CUDA

GPGPU API are very similar. This is the way to have the thread index in CUDA: In the kernel, the thread Id, the block Id and the blockDim allow the user to compute the unique thread id.

```
int index = blockIdx * blockDim + threadIdx;
```

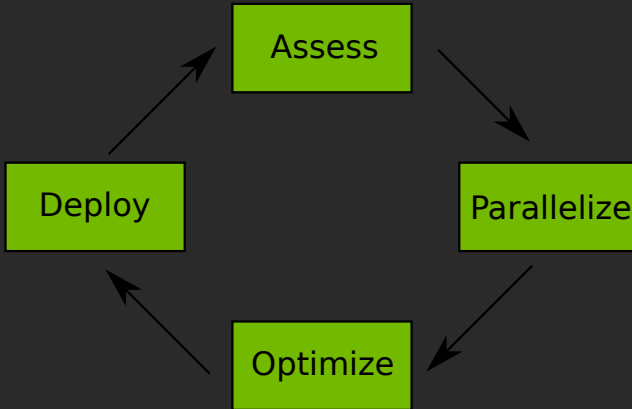
If data is stored into 2D or 3D array, it is possible to launch the kernel using a 3d vector instead of an integer and the index becomes:

```
int x = blockIdx.x * blockDim.x + threadIdx.x;  
int y = blockIdx.y * blockDim.y + threadIdx.y;  
int z = blockIdx.z * blockDim.z + threadIdx.z;
```

GPGPU usage in the industry

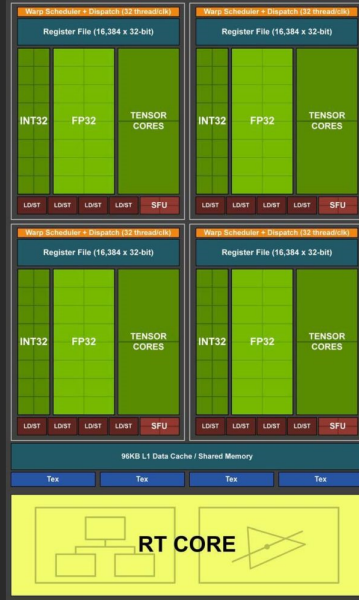
APOD

The Assess, Parallelize, Optimize, Deploy (APOD) design cycle's goal is to identify and correct bottlenecks into the application.

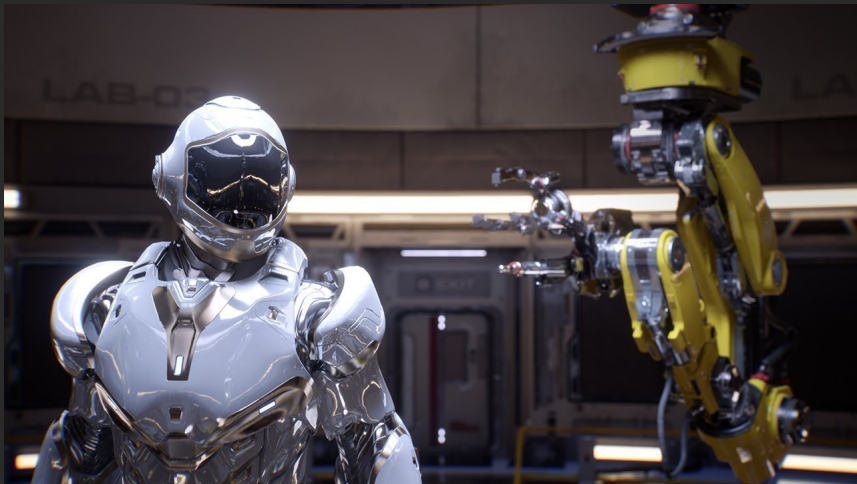


New devices

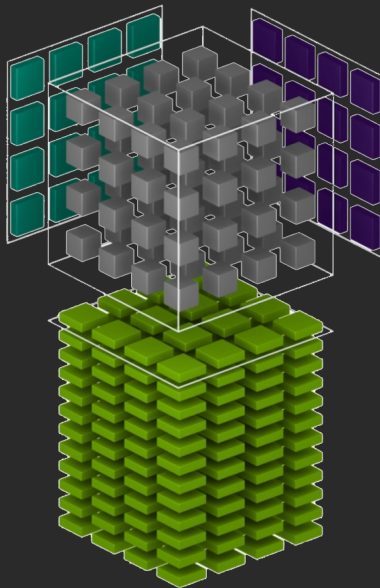
SM



RT cores



tensor cores



Domain Specific

- ▶ Rendering
- ▶ Deep Learning
- ▶ Linear Algebra and Math: Solver, Random function, Finite element method, etc...
- ▶ Signal
- ▶ Image and video
- ▶ Data oriented algorithm

Q&A